

```

import pandas as pd import numpy as np import matplotlib.pyplot as plt import
seaborn as sns from sklearn.model_selection import train_test_split, GridSearchCV from
sklearn.preprocessing import StandardScaler, PolynomialFeatures from sklearn.linear_model import
LinearRegression, Ridge, Lasso from sklearn.tree import DecisionTreeRegressor from sklearn.ensemble
import RandomForestRegressor, GradientBoostingRegressor from xgboost import XGBRegressor from
sklearn.svm import SVR from sklearn.metrics import mean_absolute_error, mean_squared_error,
r2_score from sklearn.pipeline import make_pipeline print("\n===
Loading Data ===") # Load dataset (replace with your dataset)

url = "https://raw.githubusercontent.com/ageron/handsonml2/master/datasets/housing/housing.csv"
data = pd.read_csv(url) print(f"\nData Shape: {data.shape}")

print("\nFirst 5 Rows:") print(data.head()) # Basic EDA Visualizations

plt.figure(figsize=(15, 10)) # Distribution of house prices plt.subplot(2, 2, 1)
sns.histplot(data['median_house_value'], kde=True, bins=30) plt.title('House Price Distribution')

# Correlation heatmap plt.subplot(2, 2, 2)

# Select only numeric columns numeric_data = data.select_dtypes(include=['number'])

# Compute correlation matrix corr = numeric_data.corr() # Plot heatmap sns.heatmap(corr, annot=True,
cmap='coolwarm', fmt=".1f")

plt.title('Feature Correlation') # Price vs. median income plt.subplot(2, 2, 3)

sns.scatterplot(x='median_income', y='median_house_value', data=data, alpha=0.3) plt.title('Price vs.
Income') # Price by ocean proximity

plt.subplot(2, 2, 4)

sns.boxplot(x='ocean_proximity', y='median_house_value', data=data) plt.xticks(rotation=45)
plt.title('Price by Location') plt.tight_layout() plt.show() print("\n=== Preprocessing Data ===")

# Handle missing values data.fillna(data.select_dtypes(include='number').median(), inplace=True)

# Feature engineering data['rooms_per_household'] = data['total_rooms']/data['households']
data['bedrooms_per_room'] = data['total_bedrooms']/data['total_rooms'] # Convert categorical to
numerical data = pd.get_dummies(data, columns=['ocean_proximity'])

# Select features and target

X = data.drop('median_house_value', axis=1) y =

```

```

= data['median_house_value'] # Train-test split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42) # Feature scaling
scaler

= StandardScaler()

X_train_scaled = scaler.fit_transform(X_train) X_test_scaled = scaler.transform(X_test)
print("\n=== Training Models ===") models =

{

"LinearRegression": LinearRegression(), "Ridge Regression": Ridge(alpha=1.0), "Lasso Regression":
Lasso(alpha=0.1),

"DecisionTree": DecisionTreeRegressor(max_depth=5),

"Random Forest": RandomForestRegressor(n_estimators=100, random_state=42),
"GradientBoosting": GradientBoostingRegressor(n_estimators=100, random_state=42), "XGBoost":
XGBRegressor(n_estimators=100, random_state=42),

"SVR": SVR(kernel='rbf')

} results = {} for name, model in models.items():

print(f"Training {name}...") model.fit(X_train_scaled[:1000], y_train[:1000]) y_pred

= model.predict(X_test_scaled) results[name] = {

"MAE": mean_absolute_error(y_test, y_pred), "RMSE": np.sqrt(mean_squared_error(y_test, y_pred)),

"R2": r2_score(y_test, y_pred)

}

# Display results results_df = pd.DataFrame(results).T print("\n=== Model Performance ===")
print(results_df.sort_values(by='RMSE')) print("\n=== Optimizing Best Model

===")

# Let's optimize Random Forest as it typically performs well
from sklearn.model_selection import RandomizedSearchCV # Smaller parameter grid or use
RandomizedSearchCV param_dist = {

'n_estimators': [50, 100, 200],

```

```

'max_depth':[None,10,20],

'min_samples_split':[2,5,10]

} rf = RandomForestRegressor(random_state=42)

random_search=RandomizedSearchCV(rf,param_distributions=param_dist,n_iter=5, cv=2,
scoring='neg_mean_squared_error', n_jobs=-
1,verbose=1,random_state=42,error_score='raise')random_search.fit(X_train_scaled,y_train)
best_model = random_search.best_estimator_

# Evaluate optimized model y_pred = best_model.predict(X_test_scaled)print("\nOptimized Model
Performance:") print(f"MAE:

{mean_absolute_error(y_test, y_pred):.2f}") print(f"RMSE: {np.sqrt(mean_squared_error(y_test,
y_pred)):.2f}") print(f"R2 Score: {r2_score(y_test,

y_pred):.4f}")print("\n===GeneratingVisualizations
===")

#FeatureImportanceplt.figure(figsize=(10, 6))

importances = best_model.feature_importances_ features = X.columns indices =
np.argsort(importances)[-10:]# Top 10 features plt.title('Feature Importances')
plt.barh(range(len(indices)), importances[indices], color='b', align='center') plt.yticks(range(len(indices)),
[features[i] for i in indices]) plt.xlabel('Relative Importance') plt.show() # Actual vs Predicted
plt.figure(figsize=(10, 6)) plt.scatter(y_test, y_pred, alpha=0.3)
plt.plot([y_test.min(),y_test.max()], [y_test.min(),y_test.max()], 'k--', lw=2) plt.xlabel('Actual Prices')
plt.ylabel('Predicted Prices') plt.title('Actual vs Predicted House Prices') plt.show() # Residual Plot
residuals = y_test - y_pred plt.figure(figsize=(10, 6)) plt.scatter(y_pred, residuals, alpha=0.3)
plt.axhline(y=0, color='r', linestyle='--') plt.xlabel('Predicted Prices') plt.ylabel('Residuals')
plt.title('Residual Plot')

plt.show()print("\n===ProgramExecution Complete ===")

```

OUTPUT:

